

Install modelstat

One command. About 30 seconds. No account signup, no Node, no Docker.

modelstat watches the local logs your AI tools already write (Claude Code, Cursor, and friends), summarises each session **on your machine**, and streams the cleaned result to your dashboard so you can see tokens, cost, and what you actually worked on. Prompts, responses, file contents, and secrets never leave your machine in readable form.

Contents

- [Quick install](#)
- [What happens when you run it](#)
- [Requirements](#)
- [Install per operating system](#)
 - [macOS](#)
 - [Linux](#)
 - [Windows](#)
 - [Servers & headless machines](#)
- [Summariser modes](#)
 - [Cloud \(default\)](#)
 - [Local](#)
 - [Self-hosted](#)
 - [Choosing a mode during install](#)
 - [Changing mode later](#)
- [Installer flags](#)
- [Everyday commands](#)
- [Run your own summariser engine](#)
- [Updating](#)
- [Uninstalling](#)
- [Data & privacy](#)
- [Where things live on disk](#)
- [Troubleshooting & FAQ](#)

Quick install

macOS / Linux

```
curl -fsSL https://install.modelstat.ai | sh
```

Windows (PowerShell)

```
irm https://install.modelstat.ai/ps | iex
```

That is the whole thing. It downloads a single small static binary, verifies it, pairs your device, and opens your dashboard. There is nothing else to configure to get started — the default (cloud summarising) needs no extra download.

What happens when you run it

The installer prints every step as it goes. In order:

1. **Downloads the binaries.** Two static executables — the collector (`modelstat`) and the summariser engine (`modelstat-summarizer`) — are pulled from GitHub Releases and their SHA-256 checksums are verified before anything runs. No Node.js, no Python, no package manager.
2. **Replaces any older install.** If a previous modelstat daemon (including the old Node version) is running, it is stopped first and cleanly replaced. Your device pairing in `~/.modelstat` is kept, so you stay the **same device** on the dashboard.
3. **Registers your device** with modelstat.ai and gets a claim link.
4. **(macOS) installs a menu-bar tray icon.**
5. **Asks where sessions should be summarised** — cloud, local, or self-hosted (see [Summariser modes](#)). This is a one-time consent choice. Redaction of secrets and names always runs on your machine first, in every mode.
6. **Prepares the on-device redactor** — a ~250 MB privacy model that finds and removes names/PII locally. Downloads on first run in every mode.
7. **(local mode only) downloads the summariser model** — a ~2.7 GB model (Qwen3.5-4B) so summarising can happen entirely on your machine.
8. **Installs a background service** (launchd on macOS, systemd `--user` on Linux, a Scheduled Task on Windows) so the daemon keeps running and survives reboots.
9. **Enables the Claude Code statusline** — live tokens · cost · taxonomy in your terminal (skip with `MODELSTAT_NO_STATUSLINE=1`).
10. **Detects your AI tools** and wires the modelstat MCP into them, so you can ask your assistant about your own spend.
11. **Opens your dashboard** in the browser to claim the device.

Your data starts appearing on the dashboard within seconds.

Requirements

OS	Supported
macOS	13 (Ventura) or newer — Apple Silicon or Intel
Linux	glibc, x86_64 or arm64, with <code>systemd --user</code> support
Windows	10 or 11 — x86_64 or arm64

- **No runtime dependencies.** The binaries are self-contained. You do **not** need Node, Python, or anything else installed. (This is a change from the old installer, which bootstrapped Node.)
- **No root / admin** for a normal install — everything goes under your home directory and runs as a per-user service. Use `--system` only if you want a machine-wide service (that one needs admin).
- **Disk:** ~50 MB for the binaries + ~250 MB for the on-device redactor. Add ~2.7 GB only if you choose **local** summarising.

Install per operating system

macOS

```
curl -fsSL https://install.modelstat.ai | sh
```

- Installs a **launchd** user agent, so the daemon starts at login.
- Adds a **menu-bar tray** icon for status at a glance.
- Everything lives under `~/.modelstat`. No `sudo` required.

Linux

```
curl -fsSL https://install.modelstat.ai | sh
```

- Installs a **systemd --user** service (`modelstat.service`).
- Requires user lingering to run when you are not logged in: `loginctl enable-linger "$USER"`.
- Everything lives under `~/.modelstat`. No `sudo` required.

Windows

```
irm https://install.modelstat.ai/ps | iex
```

- Installs a **Scheduled Task** named `modelstat` that runs at logon.
- Everything lives under `%USERPROFILE%\modelstat`. No admin required.

Pass flags after the piped command like this:

```
& ([scriptblock]::Create((irm https://install.modelstat.ai/ps))) -Mode local
```

Servers & headless machines

On a machine with no browser / no interactive terminal, tell the installer everything up front so it never has to prompt:

```
curl -fsSL https://install.modelstat.ai | sh -s -- --mode cloud --yes --no-browser
```

- `--mode` is **required** on a fresh non-interactive install (there is no silent default — you must state your consent choice once).
- `--yes` accepts the defaults; `--no-browser` skips opening the dashboard.
- Add `--system` to install a machine-wide service (needs root/admin).

Summariser modes

Every mode runs **redaction on your machine first** — secrets and names/PII are stripped locally before anything is sent anywhere. What differs is only *where the summary itself is produced*.

Cloud (default)

modelstat's servers summarise your cleaned, redacted turns.

- **Resource:** no local model; negligible RAM, CPU, and battery on your machine.
- **Privacy:** your cleaned, redacted turns are uploaded and summarised server-side.
- **Best for:** most people. Nothing extra to install; lowest local footprint.

Local

A bundled model summarises entirely on **this** machine.

- **Resource:** ⚠ downloads a ~2.7 GB model (Qwen3.5-4B) and uses ~4 GB RAM plus extra battery/CPU while summarising.
- **Privacy:** only a short (≤ 240 -character) abstract is uploaded — the raw turns never leave your machine.
- **Best for:** people who want the strongest privacy and have the RAM/disk to spare.

Self-hosted

Your organisation's own summariser engine, at a URL you provide.

- **Resource:** nothing extra here — summarising runs on the engine URL you point at.
- **Privacy:** only the abstract reaches modelstat; the cleaned excerpts go to your org's engine, not ours.

- **Best for:** teams that want to keep summarisation inside their own infrastructure. See [Run your own summariser engine](#).

Choosing a mode during install

```
# Cloud (default – you can also just omit --mode when interactive)
curl -fsSL https://install.modelstat.ai | sh -s -- --mode cloud

# Local
curl -fsSL https://install.modelstat.ai | sh -s -- --mode local

# Self-hosted (point at your org's engine)
curl -fsSL https://install.modelstat.ai | sh -s -- --mode self-hosted --url
http://llm.acme.internal:4321
```

On Windows use `-Mode` / `-Url` .

Changing mode later

```
modelstat mode          # show the current mode
modelstat mode local     # switch to local (downloads the model, arms the engine)
modelstat mode cloud     # switch back to cloud (removes the local engine, keeps model
files)
modelstat mode self-hosted --url http://llm.acme.internal:4321
```

Switching mode reconfigures the local engine, refreshes the background service, and takes effect on the next scan. You can also pin a mode with the environment variable `MODELSTAT_SUMMARIZER_MODE` , or a self-hosted URL with `MODELSTAT_SUMMARIZER_URL` ; when set, these override the stored choice.

Installer flags

Pass flags after `| sh -s --` on macOS/Linux, or directly on Windows.

Flag (Unix)	Flag (Windows)	Meaning
<code>--component daemon summarizer</code>	<code>-Component</code>	Install the collector (default) or a standalone summariser engine.
<code>--user (default) / --system</code>	<code>-System</code>	Per-user install, or a machine-wide service (needs root/admin).
<code>--version X.Y.Z</code>	<code>-Version</code>	Install a specific release instead of the latest.
<code>--mode cloud local self-hosted</code>	<code>-Mode</code>	Summariser mode (required on a fresh non-interactive install).
<code>--url <URL></code>	<code>-Url</code>	Engine URL for <code>self-hosted</code> mode.
<code>--yes / -y</code>	<code>-Yes</code>	Accept defaults; don't prompt.
<code>--no-browser</code>	<code>-NoBrowser</code>	Don't open the dashboard at the end.
<code>--no-auto-update</code>	<code>-NoAutoUpdate</code>	Don't let the daemon update itself automatically.

Auto-update is on by default (it is how you receive fixes). It is turned off automatically if you install a **pre-release** build, or if you pass `--no-auto-update`. Toggle it any time with `modelstat autoupdate on|off`.

Everyday commands

Run these in any terminal after install:

Command	What it does
<code>modelstat</code>	Re-run onboarding (safe to repeat — idempotent). Aliases: <code>connect</code> , <code>reinstall</code> .
<code>modelstat status</code>	Pairing, service state, and a live usage snapshot. Add <code>--json</code> for machines.
<code>modelstat jobs</code>	The local pipeline queue and recent activity.
<code>modelstat mode</code>	Show or change where sessions summarise.
<code>modelstat discover</code>	List the AI tools and signed-in accounts detected locally (read-only).
<code>modelstat sync --session <id> [-wait]</code>	Force one session to be scanned right now.
<code>modelstat reset</code>	Re-read and re-summarise everything from scratch (wipes local cursors).
<code>modelstat upgrade</code>	Update to the latest release now.
<code>modelstat autoupdate on off toggle</code>	Set the auto-update preference.
<code>modelstat stop</code>	Remove the background service (your pairing is kept). Aliases: <code>remove</code> , <code>uninstall</code> .
<code>modelstat paths</code>	Show resolved file paths (state, identity, logs, API).
<code>modelstat token</code>	Print this device's bearer token (treat it like a password).

Run your own summariser engine (self-hosted)

For a team that wants summarisation to happen on its **own** box, install the engine component on a server and point your collectors at it.

On the engine server:

```
curl -fsSL https://install.modelstat.ai | sh -s -- --component summarizer
```

This runs `modelstat-summarizer setup`, which:

- asks for a **bind address** and **port** (binding to `0.0.0.0` requires typing `expose` to confirm — the engine has **no authentication**, so only do that behind a firewall or reverse proxy you trust),
- downloads the model,
- installs the engine as a background service,
- asks whether to enable daily auto-update (default: off, so a shared box changes predictably),
- and prints the exact command your collectors should run.

On each collector machine, point it at the engine:

```
modelstat mode self-hosted --url http://<engine-host>:<port>
```

Engine management commands (on the server):

Command	What it does
<code>modelstat-summarizer status</code>	Show config and probe the running engine.
<code>modelstat-summarizer serve</code>	Run the engine in the foreground.
<code>modelstat-summarizer stop</code>	Stop the engine service (keeps it installed).
<code>modelstat-summarizer uninstall</code>	Remove the engine service (keeps model files).
<code>modelstat-summarizer upgrade</code>	Update the engine to the latest release.

Updating

- **Automatic (default):** the daemon checks for new releases and updates itself in the background. Nothing to do.
- **Manual:** `modelstat upgrade`.
- **Turn auto-update off/on:** `modelstat autoupdate off` / `modelstat autoupdate on`.

Updates are verified by SHA-256 (and a signature when signing is enabled), applied by swapping the binaries, health-checked, and rolled back automatically if the new version fails to come up.

Uninstalling

```
modelstat stop          # or: modelstat remove / modelstat uninstall
```

This removes the background service, the menu-bar tray, and the Claude Code statusline. Your device pairing is **kept** in `~/.modelstat`, so `modelstat` will re-enable everything later.

To remove it completely, also delete the home directory:

```
rm -rf ~/.modelstat      # macOS / Linux
```

```
Remove-Item -Recurse -Force "$env:USERPROFILE\.modelstat"  # Windows
```

To reclaim disk without fully uninstalling, delete just the models: `rm -rf ~/.modelstat/models/`.

Data & privacy

- **Redaction is always local.** Secrets and names/PII are removed on your machine before anything is sent, in every mode.
- **Prompts, responses, file contents, and secrets stay on your machine.** In cloud mode, only *cleaned, redacted* turns are uploaded. In local and self-hosted mode, only a short abstract leaves your machine.
- **Read-only access.** The collector only reads the log directories your AI tools already write to. It does not need root/sudo for a normal install.

Where things live on disk

Everything is under one directory:

- **macOS / Linux:** `~/.modelstat`
- **Windows:** `%USERPROFILE%\modelstat`
- **--system installs:** `/var/lib/modelstat` (Linux) or `%ProgramData%\modelstat` (Windows)

Inside it:

Path	Contents
<code>bin/</code>	the <code>modelstat</code> and <code>modelstat-summarizer</code> binaries
<code>models/</code>	downloaded models (redactor, and Qwen for local mode)
<code>identity.json</code>	your device pairing (keep this to stay the same device)
<code>state.json</code>	scan cursors and local state
<code>logs/</code>	daemon logs
<code>summarizer.json</code>	engine config (self-hosted / local engine only)

Troubleshooting & FAQ

Nothing shows up on the dashboard. Run `modelstat status` to check pairing and the service, and `modelstat jobs` to see the local queue. If the device is not claimed, open the claim link the installer printed.

I want to re-process my history from scratch. `modelstat reset` wipes the local cursors so the next scan re-reads and re-summarises everything.

Does it need Node / Python / Docker? No. The binaries are self-contained.

Does it run as root? No, not for a normal install — it runs as a per-user service under your home directory. `--system` is the only mode that installs a machine-wide service.

How do I install a specific version? Pass `--version X.Y.Z` (Unix) or `-Version X.Y.Z` (Windows).

Is my device the same after re-installing or upgrading? Yes. As long as `~/.modelstat` (specifically `identity.json`) is intact, you remain the same device on the dashboard — no duplicate.

Help / support: <https://modelstat.ai/install>